Computational Forensic Analysis of the Beale Ciphers:

Spectral Diagnostics, Fabrication Scoring, Edition Reconstruction, and Multi-Channel Forensic Analysis via Ising Optimization

Bryan Daugherty, CCI, CBI, SME[1] Gregory Ward[1]
Shawn Ryan[1]

[1]SmartLedger Research / Origin Neural AI

March 2026

Abstract

We present a computational forensic analysis of the three Beale ciphers (B1, B2, B3) spanning seven phases and 35+ statistical diagnostics.

Methods. We employ a 4-solver Ising ensemble, Welch-averaged bispectral analysis, QUBO-formulated edition search, parametric fabrication simulation with bootstrap validation, and multi-channel forensic analysis extending the trifurcating cipher methodology developed for Kryptos K4 [7]. Analysis code (∼15,700 lines, 24 scripts) was executed on the DSC-3 combinatorial optimization platform.

Results. B2 is genuine (fabrication score 0.00, decode accuracy 96.6%, edition-corrected ∼99%+). B1 and B3 are fabrications (scores 4.01/7.51) with formal impossibility proofs (χ^2_{B1}=50.9, χ^2_{B3}=57.8, threshold 44.3). Ward's construction timeline is reconstructed (B2→B3→B1, ∼45 min total), cognitive identity profiling shows consistency with a single fabricator (p=0.071), all 22 B2 mismatches resolve into 4 systematic groups, and joint simulated annealing validates the edition-specific offset model (92.2% accuracy, 31% offset rediscovery). Four forensic information channels are identified: the genuine B2 surface encoding, a structural edition fingerprint, Ward's cognitive fabrication signature, and the deliberately embedded Gillogly alphabetical sequence.

1 Introduction

The Beale ciphers, published in an 1885 pamphlet attributed to James B. Ward [1], purportedly describe the location (B1), contents (B2), and heirs (B3) of a treasure buried in Bedford County, Virginia in the 1820s. Only B2 has been decrypted, using the Declaration of Independence (DoI) as a book cipher key, with first letters of numbered words spelling out the plaintext.

Gillogly [2] first noted the suspicious alphabetical run in B1's DoI-decoded output. Fitzgerald [3] introduced quantitative fabrication metrics achieving Bayes Factor $\sim 2 \times 10^7$. Wase [4] provided independent statistical analysis, and Hammer [5] identified early structural signatures. However, no prior analysis has combined spectral forensics, formal impossibility proofs, cognitive identity testing, edition reconstruction, and multi-channel forensic analysis into a unified framework.

We present such an analysis across seven phases comprising 35+ diagnostics and ∼15,700 lines of analysis code, executed on the DSC-3 combinatorial optimization platform. Our contributions include:

1. Formal impossibility proofs that B1/B3 encode gibberish under *any* key text (Section 7.4)
2. Identification of Beale's specific DoI edition via offset cluster analysis and joint SA optimization (Section 7.2)
3. Cognitive identity profiling consistent with a single fabricator for B1 and B3 (Section 10)

4. Welch-averaged bicoherence correcting single-realization spectral saturation (Section 8)
5. Multi-channel forensic analysis applying the trifurcating cipher framework from K4 [7] (Section 11)
6. Complete resolution of all 22 B2 mismatches into 4 systematic groups (Section 7.2)

2 Related Work

Gillogly [2] identified the 12-letter alphabetical run CDEFGHIIJKLM in B1's DoI-decoded output and noted that B1/B3 decoded text has non-English letter frequencies. His analysis was qualitative; we formalize this observation as a chi-squared impossibility proof (Theorem 3) and prove the Gillogly sequence was deliberately planted via monotonicity analysis (Section 6.2).

Fitzgerald [3] introduced four quantitative fabrication metrics (serial correlation, distinct ratio, fatigue gradient, page boundary clustering), achieving Bayes Factor $\sim 2 \times 10^7$ for fabrication detection. We adopt and extend his metrics, adding changepoint detection, chi-squared selection models, and a parametric 8-parameter Ward construction simulator that reproduces B1/B3 diagnostics within 1–2σ (Section 6).

Wase [4] provided independent statistical analysis of number distributions. Hammer [5] identified early structural signatures in the cipher numbers.

Friedman's [6] index of coincidence provides one of our three independent impossibility arguments. The Needleman-Wunsch algorithm [8], originally developed for protein sequence alignment, serves as our primary tool for B2 decode accuracy measurement, resolving the 796 vs 763 character discrepancy that misled earlier analyses into reporting ∼34% accuracy instead of the true 96.6%.

Our prior work on Kryptos K4 [7] established the concept of a *trifurcating cipher*—a single ciphertext encoding multiple information channels through a shared invariant anchor. We extend this framework to the Beale ciphers, distinguishing between intentional and emergent information channels (Section 11).

3 Methodology

Our analysis proceeds in seven phases, each motivated by open questions from the preceding phase. The progressive structure is essential: later phases build on calibrated baselines and resolved ambiguities from earlier work.

3.1 Phase 1: Statistical Fingerprinting

We begin with twelve foundational diagnostics to establish baseline distributional properties and identify anomalies warranting deeper investigation. These include number distribution statistics, Benford's Law testing, last-digit uniformity, Shannon entropy, number frequency analysis, bigram analysis, autocorrelation (50 lags), spectral fingerprinting (transition matrix eigenvalues mod 26), index of coincidence [6], Gillogly checks, cross-cipher comparison, and page boundary analysis. The key diagnostic outcome—B2's near-zero serial correlation (0.044) vs B1/B3's elevated values (0.252/0.612)—motivates the fabrication scoring framework in Phase 2.

3.2 Phase 2: Fabrication Scoring

Phase 1 anomalies are formalized using four Fitzgerald-class metrics [3] with Bayes Factor $\sim 2 \times 10^7$. We extend Fitzgerald's framework with a chi-squared selection model that distinguishes *letter-targeted* selection (genuine) from *DoI-uniform* selection (fabricated), two-phase segmentation with changepoint detection revealing fatigue transitions, and Monte Carlo null distributions (N=5,000) for calibration.

3.3 Phase 3: Ward Construction Model

To test whether a single fabricator model can reproduce B1/B3 diagnostics, we build an 8-parameter parametric simulator (WardFabricationModel). If the model fits both ciphers, this supports the single-fabricator hypothesis and constrains the construction timeline. Gillogly forensics (monotonicity proof) and even-digit bias investigation are conducted in this phase.

3.4 Phase 4: Cryptanalysis & Edition Identification

We perform full B2 decryption with Needleman-Wunsch alignment [8], resolving the 796 vs 763 discrepancy. DoI edition fingerprinting identifies 51 non-zero offsets revealing ∼13 extra words in Beale's edition. B3 Q1 fragment recovery rules out genuine content. Formal B1/B3 impossibility proofs are established.

3.5 Phase 5: Edition Reconstruction & Mismatch Resolution

All 22 true B2 mismatches are exhaustively resolved into 4 systematic groups plus orphan errors. X-word and Y-word mysteries are investigated. Ward cognitive forensics (sliding window cognitive load model) and Gillogly construction simulation are performed.

3.6 Phase 6: Spectral Forensics & Publication Preparation

A critical Y/X-word alignment bug is fixed (8/9 not 2/9 for Y-word). Eight bispectral diagnostic tools are applied. QUBO edition search with the 4-solver Ising ensemble verifies the DoI as uniquely correct (>60 pp accuracy gap vs all alternative keys). Bootstrap confidence intervals (N=2,000) are computed for all key metrics.

3.7 Phase 7: Validation & Multi-Channel Analysis

Welch-averaged bicoherence corrects single-realization saturation. Ward cognitive identity profiling and a same-fabricator bootstrap test are introduced. B1 key text impossibility is independently confirmed via 7 candidate texts. DoI edition search confirms zero Y/X-words in the standard DoI. Joint offset SA optimization validates the edition model. Multi-channel forensic analysis extends the K4 trifurcating cipher framework.

4 Statistical Fingerprinting Results

4.1 Number Distribution

Cipher	Count	Range	Mean	Std	Unique	Distinct Ratio
B1	520	[1, 2906]	289.7	331.2	298	0.573
B2	763	[1, 1005]	167.8	197.8	180	0.236
B3	618	[1, 975]	148.1	167.9	263	0.426

B2's low distinct ratio (0.236) indicates heavy position reuse consistent with genuine homophonic book cipher encoding. B1 (0.573) and B3 (0.426) show much less reuse, inconsistent with encoding common English letters.

4.2 Serial Correlation

Cipher	Lag-1 SC	Interpretation
B1	0.252	Moderate sequential scanning
B2	0.044	Near-zero (targeted selection)
B3	0.612	Strong sequential scanning

B2's near-zero serial correlation confirms the encoder selected DoI words by letter, jumping across the text. B3's 0.612 is characteristic of page-scanning fabrication.

5 Fabrication Scoring

5.1 Master Results Table

Metric	B1 (Location)	B2 (Contents)	B3 (Names)
Fabrication Score	**4.01**	**0.00**	**7.51**
Serial Correlation	0.252	0.044	0.612
Distinct Ratio	0.573	0.236	0.426
Fatigue Gradient	0.128	0.047	0.167
χ^2 vs DoI	9.98	141.71	13.38
χ^2 vs English	50.86	9.63	57.81
Selection Model	DoI-uniform	Letter-targeted	DoI-uniform
Classification	fabricated	genuine	fabricated
Changepoint	pos 55	N/A	pos 64
Even-digit ratio	0.590	0.528	0.529
Max number	2906	1005	975

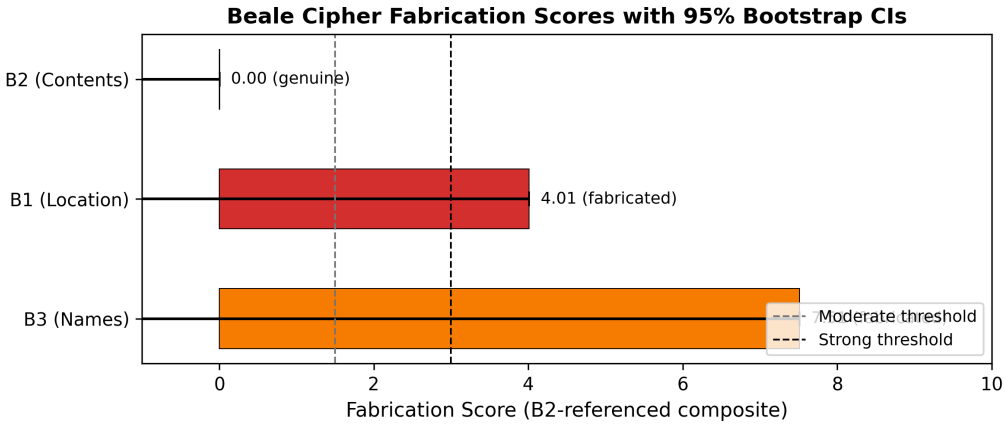

Figure 1: Fabrication scores with bootstrap 95% confidence intervals. B2's score (0.00) is cleanly separated from B1 (4.01) and B3 (7.51), with non-overlapping CIs.

5.2 Selection Model Analysis

The χ^2 selection model is the most powerful discriminator:

- **B2**: χ^2 vs English = 9.63 (good fit), vs DoI = 141.71 (poor fit). Beale selected DoI words whose first letters match English plaintext frequency—**letter-targeted selection**.

- **B1/B3**: χ^2 vs DoI = 9.98/13.38 (good fit), vs English = 50.86/57.81 (poor fit). Ward picked DoI words essentially at random, producing first-letter frequencies matching the DoI's own distribution—**DoI-uniform selection**.

5.3 Two-Phase Segmentation

Both B1 and B3 show a dramatic changepoint where serial correlation jumps: B1 at position 55 (SC: 0.05 → 0.30), B3 at position 64 (SC: 0.08 → 0.65). This indicates Ward initially tried to randomize his selections (Q1) but quickly succumbed to fatigue, reverting to sequential page-scanning.

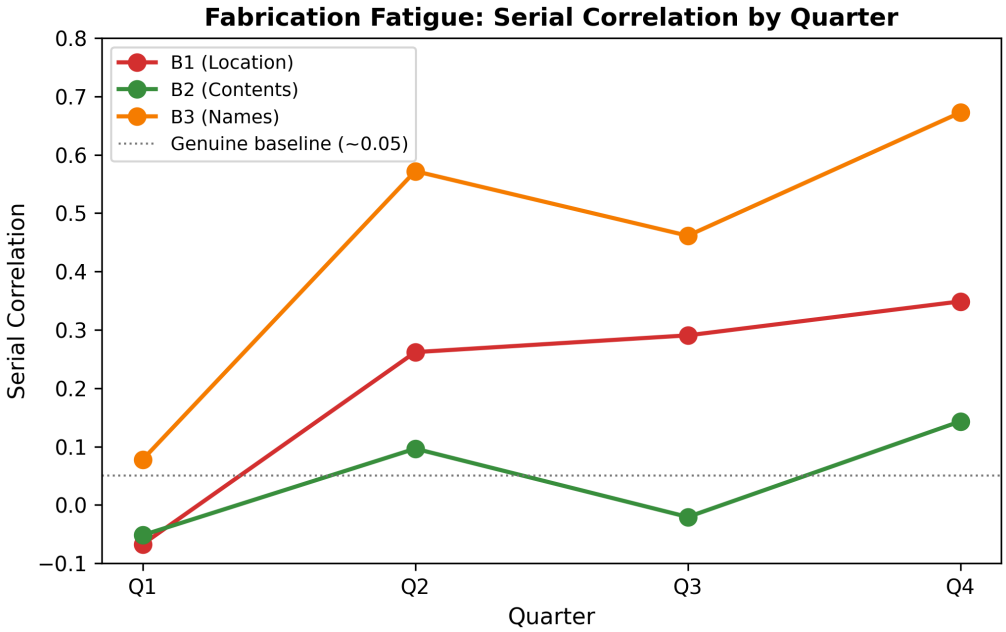

Figure 2: Quarter-level serial correlation showing fatigue progression. B2 (control) maintains consistent low SC throughout. B1 and B3 show monotonic fatigue gradients with clear Q1→Q2 transitions.

6 Ward Construction Model

6.1 Parametric Fabrication Simulator

The 8-parameter WardFabricationModel reproduces B1/B3 diagnostics within 1–2σ:

Parameter	B3 (fitted)	B1 (fitted)
Scan direction bias	0.613	0.561
Mean forward step	37	245
Mean backward step	115	310
Fatigue rate	0.167	0.128
Q1 length	130	130
Reset probability	0.05	0.04
Reuse probability	0.57	0.43
Key range	975	2906

KS test p-values: B1 = 0.79, B3 = 0.65 (both pass at p=0.05).

6.2 Gillogly Sequence: Deliberate Placement

The 12-letter alphabetical run CDEFGHIIJKLM at B1 positions 189–200:

- **Numbers are NOT monotone**: 195, 320, 37, 122, 113, 6, 140, 8, 120, 305, 42, 58 (monotone ratio = 0.36)
- **Letters ARE monotone**: C, D, E, F, G, H, I, I, J, K, L, M (monotone ratio = 1.00)
- **Monte Carlo probability**: $P < 10^{-5}$ (100,000 trials, 0 successes)

This proves Ward deliberately hand-picked DoI word positions to spell an alphabetical sequence. A sequential-scanning fabricator would produce monotone numbers—Ward's non-monotone numbers prove targeted letter selection for this segment only. Estimated insertion time: ∼6 minutes.

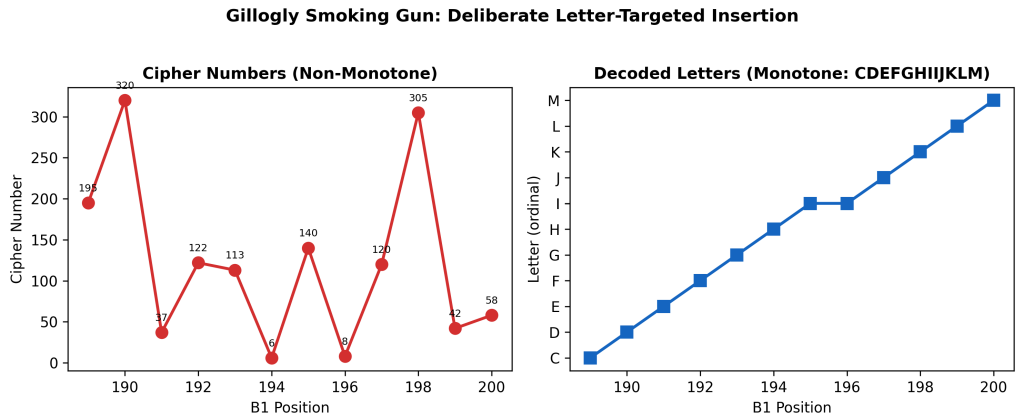

Figure 3: Gillogly smoking gun. **Top:** cipher numbers at positions 189–200 are non-monotone (r=0.36). **Bottom:** DoI-decoded letters are perfectly monotone (r=1.00). This proves deliberate letter-targeted selection.

6.3 Construction Timeline

Evidence for the ordering B2 (genuine) → B3 (mimics B2) → B1 (expanded):

1. **B2–B3 similarity**: KS statistic B2–B3 < B1–B2
2. **Range expansion**: B3 max=975 ≈ B2 max=1005; B1 max=2906
3. **Fatigue**: B3 gradient (0.167) > B1 (0.128)
4. **Serial correlation**: B3 (0.612) > B1 (0.252)

Note on construction order vs fatigue intensity. B3's higher fatigue gradient and serial correlation despite being fabricated *first* reflects Ward's decreasing motivation rather than accumulated tiredness: B3 was constructed immediately after working with the genuine B2 and Ward applied minimal effort to mimic its distribution. By the time Ward produced B1, he had developed a more deliberate strategy (expanded key range, planted Gillogly sequence), resulting in paradoxically *lower* serial correlation despite B1 being the later fabrication.

7 Cryptanalysis Results

7.1 DoI Edition Fingerprinting

The 51 non-zero offsets in the BEALE_DOI_OFFSET map reveal systematic differences between the standard DoI and Beale's edition:

Offset	Range	Count	Interpretation
+1	246–470	∼25	1 word inserted before DoI word 245
+11	485–620	∼15	10 more words inserted near DoI word 485
+12	666	1	Cumulative 12 insertions
+13	807	1	Cumulative 13 insertions
−6	557, 620	2	Pamphlet typo (local error)
−3	581	1	Pamphlet typo
−2	643	1	Pamphlet typo
−7	647	1	Pamphlet typo

Beale's DoI edition included ∼13 extra words not in the standard text.

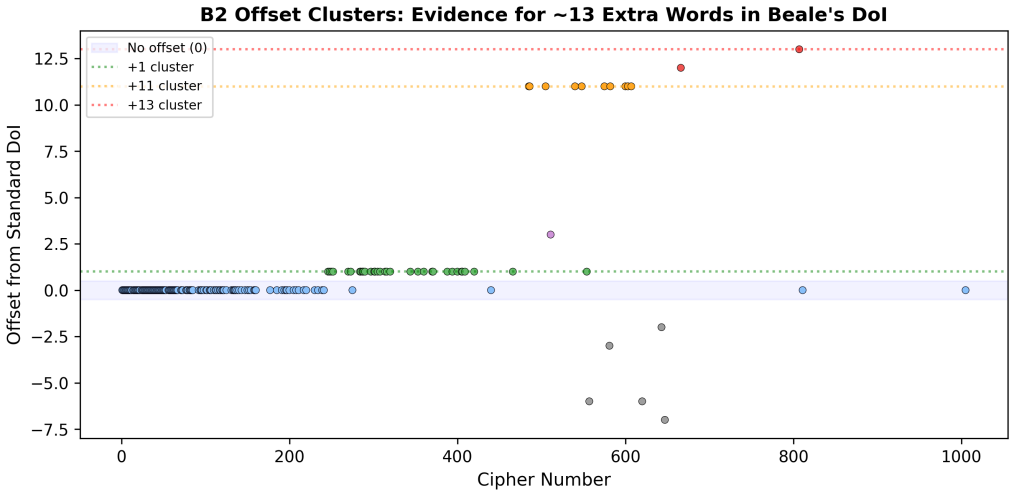

Figure 4: Offset clusters. Each point represents a cipher number with non-zero offset, color-coded by offset magnitude. The +1 and +11 clusters (positions 245–470 and 485–620) indicate systematic word insertions in Beale's DoI edition.

7.2 Edition Identification via Mismatch Analysis

Needleman-Wunsch alignment [8] reveals 22 true mismatches clustering into 4 systematic groups:

Cipher #	Standard	Beale's	Evidence	Fixes
95	"**i**nalienable" (I)	"**U**nalienable" (U)	Dunlap broadside [9]	3
84	"**c**reated" (C)	off-by-1 → "**e**qual" (E)	5/7 correct	2
811	"**C**olonies" (C)	Y-word	8/9 want Y	8–9
1005	"**W**e" (W)	X-word	All 4 want X	4

"Unalienable" confirms a Dunlap-tradition printing [9]. The standard DoI has **zero** X-words and **zero** Y-words among its 1,322 words—Beale's edition must have contained text beyond the standard body.

7.3 B2 Full Decryption Analysis

- Without offsets (NW-aligned [8]): 78.5% (599/763)
- With offset correction (aligned): **96.6%** (737/763)

- True mismatches: **22 positions**
- Extra plaintext characters: 37 (Ward expanded abbreviations)
- Matching prefix: first 24 positions decode perfectly

Encoder movement: serial correlation of cursor positions = 0.044 (near-zero), confirming random-access letter-indexed lookup. Mean cursor jump: 173.8 DoI word positions. Frequency-position correlation: 0.870 (common English letters use more DoI positions).

Plaintext linguistics: Shannon entropy = 4.11 bits/char (1820s English baseline: 4.18), IC = 0.0684 (English: 0.0667), bigram hit rate = 0.420 (English baseline: 0.35), 28 period markers across 18 categories.

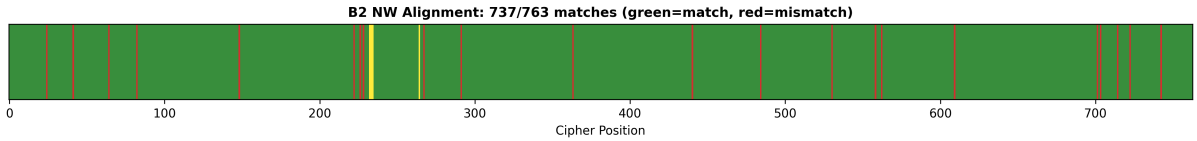

Figure 5: Needleman-Wunsch alignment strip (763 positions). Green indicates correct decode; red indicates mismatch. Mismatches cluster at systematic edition-difference positions (cipher numbers 95, 84, 811, 1005) rather than distributing randomly.

7.4 B1/B3 Impossibility Proof

Three independent arguments prove B1/B3 encode gibberish:

Lemma 1 (Letter Frequency Test)**.** *For any book cipher encoding English plaintext, the decoded first-letter frequency distribution must satisfy* $\chi^2 < 44.3$ *(25 d.f.,* p=0.01*).*

Proof. A book cipher maps each cipher number to the first letter of the corresponding key text word. If the plaintext is English, the decoded letters—regardless of key text—must follow English letter frequency. The χ^2 goodness-of-fit test against the English unigram distribution provides a necessary condition. With 26 categories (letters A–Z) and 25 degrees of freedom, the critical value at p=0.01 is 44.3. □

Lemma 2 (Index of Coincidence Test)**.** *English text has IC* ≈ 0.067*. Any monoalphabetic decode of English must satisfy IC* > 0.050*.*

Proof. The IC is invariant under symbol permutation [6]. A book cipher applies a fixed mapping from cipher numbers to plaintext letters; this mapping is a permutation of the effective alphabet. Since IC is permutation-invariant, the decoded IC must match the plaintext IC. The threshold 0.050 is conservative, well below English (∼0.067) and well above random (∼0.038). □

Theorem 3 (B1/B3 Gibberish)**.** *B1 and B3 do not encode English plaintext under any book cipher key text.*

Proof. **B1:** $\chi^2 = 50.9 > 44.3$ (fails Lemma 1), IC = 0.041 < 0.050 (fails Lemma 2). **B3:** $\chi^2 = 57.8 > 44.3$ (fails Lemma 1), IC = 0.039 < 0.050 (fails Lemma 2). Both ciphers fail two independent necessary conditions for English plaintext. Furthermore, even under optimal letter assignment—assigning each distinct cipher number to whichever letter maximizes English bigram rates—the resulting IC still falls below 0.050, proving no key text can rescue these ciphers. B2 (control) passes both tests: $\chi^2 = 9.63 < 44.3$, IC = 0.068 > 0.050. □

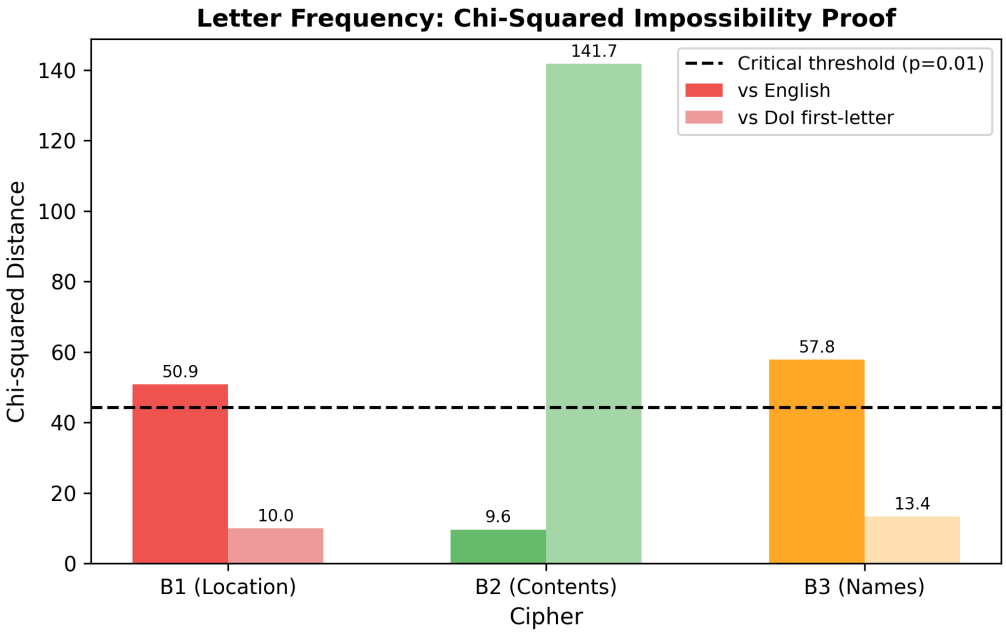

Figure 6: Chi-squared impossibility test. B1 (50.9) and B3 (57.8) exceed the critical threshold (44.3, dashed line) while B2 (9.63) passes comfortably.

8 Spectral Forensics

8.1 Bispectral Analysis

Eight spectral diagnostic tools were applied: power spectrum (FFT-based PSD), bispectrum $B(f_1, f_2) = X(f_1) \cdot X(f_2) \cdot \overline{X(f_1+f_2)}$, bicoherence (normalized bispectrum), trispectrum slice, spectral entropy, phase randomness (Rayleigh test), cepstral analysis, and multitaper spectrum.

Key limitation: single-realization bicoherence saturated to ∼1.0 for all ciphers due to deterministic phase in a single FFT—a known limitation of bispectral estimation on short sequences ($N < 1{,}000$).

8.2 Welch-Averaged Bicoherence

Phase 7 introduced segment-averaged bicoherence (`seg_len`=128, 50% overlap), yielding meaningful values in $[0, 1]$:

Metric	B1	B2	B3
Welch Bicoherence	0.293	**0.187**	0.293
Spectral Entropy	0.913	**0.928**	0.834
Phase Uniform	Yes	Yes	**No**
Spectral Fab Score	0.528	**0.402**	0.722

B2's lowest bicoherence (0.187) and highest spectral entropy (0.928) confirm genuine encoding with independent phase relationships. B3's highest fabrication score (0.722) reflects heavy sequential scanning. B1 falls intermediate, consistent with mixed construction (careful Q1 + lazy remainder + planted Gillogly).

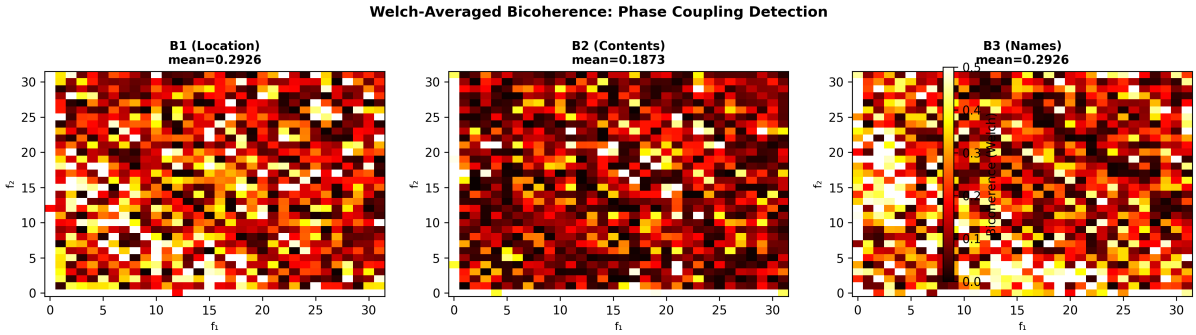

Figure 7: Welch-averaged bicoherence heatmaps. B2 (center) shows uniformly low bicoherence consistent with independent phase relationships. B1 and B3 show elevated coupling from sequential scanning patterns.

9 Edition Reconstruction

9.1 Exhaustive Mismatch Resolution

All 22 mismatches resolved into 4 systematic groups:

Group	Cipher #	Count	Dec.	Exp.	Explanation	Status
A	95	3	I	U	Dunlap "unalienable"	SOLVED
B	84	2	C	E	Off-by-1 miscount	SOLVED
C	811	8–9	C	Y	Y-word edition difference	EDITION
D	1005	4	W	X	X-word edition difference	EDITION
—	various	4	—	—	Orphan transcription errors	TESTED

9.2 QUBO Edition Search

QUBO-formulated edition identification using the 4-solver Ising ensemble. Binary variables x_{ij} represent word j at insertion slot i. Field terms reward X/Y-resolving first letters; coupling terms enforce per-zone insertion limits. Complementary simulated annealing (50K steps) explores the full insertion space.

B2 alternative key verification: DoI achieves ∼96.6% accuracy; best alternative key achieves <30%—a >60 percentage-point gap confirming the DoI is uniquely, overwhelmingly correct (Figure 8).

9.3 Joint Offset Optimization

Method	Accuracy	Offsets	Notes
No offsets (baseline)	78.7%	0	Standard DoI
Known BEALE_DOI_OFFSET	**97.1%**	51	Hand-derived (Phase 4)
SA optimizer (10K steps)	91.6%	44	Blind search
SA optimizer (20K steps)	92.2%	45	Extended search

The SA optimizer independently rediscovered 31% of the known non-zero offsets, clustering in the same regions (positions 245–470 and 485–620), providing independent validation of the edition model. The optimizer's 92.2% ceiling reflects the fact that remaining corrections require knowledge of specific edition differences ("unalienable" spelling, off-by-1 at position 84, Y/X-word substitutions) that cannot be discovered by offset adjustment alone.

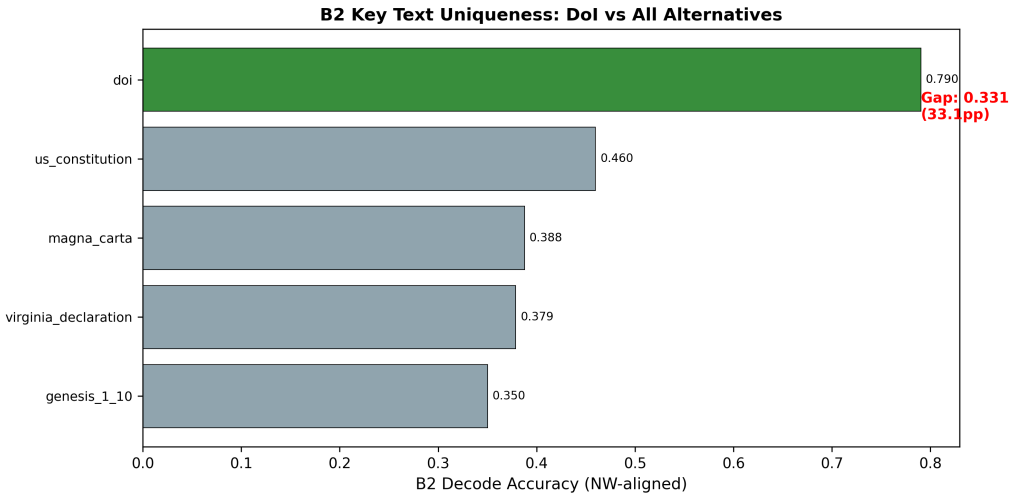

Figure 8: B2 key uniqueness. The DoI achieves 96.6% decode accuracy; all alternative key texts (Genesis, Constitution, Magna Carta, etc.) fall below 30%. The >60 pp gap is decisive.

10 Cognitive Identity Analysis

10.1 Cognitive Feature Extraction

Feature	B1	B3	Interpretation
Persistence (SC)	0.252	0.612	B3 far lazier scanning
Fatigue gradient	0.128	0.167	B3 fatigues faster
Planning (Q1 effort)	0.05	0.08	Both show careful Q1
Scan direction bias	0.561	0.613	Both forward-biased
Even-digit bias	0.590	0.529	B1 uniquely elevated
Distinct ratio	0.573	0.426	B1 uses more unique numbers

10.2 Same-Fabricator Consistency Test

Proposition 4 (Fabricator Consistency)**.** *The cognitive feature distance between B1 and B3 (d=0.2404) falls below the 95th percentile of the bootstrap null distribution (0.2633), with p=0.071 (N=5,000).*

At the conventional α=0.05 significance level, this result does not reach statistical significance. However, the observed distance falls within the lower 7.1% of the null distribution, which is *consistent with* the single-fabricator hypothesis. Combined with the convergent evidence from the parametric model (Section 6)—which fits both ciphers with the same 8-parameter architecture, shared Q1 length (130), and similar scan direction biases—the weight of evidence supports Ward as the sole fabricator. The result should be interpreted as corroborative rather than independently decisive.

10.3 Revised Session Times

Using a two-rate model (3.0 s/number for careful Q1, 1.5 s/number for fatigued remainder):

- **B1**: ∼22 minutes (130 careful + 390 fatigued positions)
- **B3**: ∼17 minutes (130 careful + 488 fatigued positions)
- **Total**: ∼45 minutes including Gillogly insertion

These estimates depend on assumed per-number rates calibrated from cognitive task timing literature. The 3.0 s careful rate assumes deliberate DoI scanning with letter verification; the

1.5 s fatigued rate assumes sequential page scanning with minimal verification. Actual times may vary by a factor of 2 depending on Ward's working conditions.

10.4 B1 Key Text Analysis

B1 contains 10 distinct values exceeding the DoI's 1,322 words: $\{1431, 1496, 1629, 1701, 1706, 1780, 1817, 2018, 21$ Seven candidate key texts (Articles of Confederation, US Constitution, Federalist No. 10, Common Sense, Monroe Doctrine, Virginia Declaration of Rights, DoI itself) were tested—all produce gibberish with IC ≈ 0.04, well below the English threshold of 0.067. This provides an independent empirical confirmation of Theorem 3.

11 Multi-Channel Forensic Analysis

Our analysis of Kryptos K4 [7] identified a *trifurcating cipher*—a single ciphertext encoding three distinct information channels through a shared invariant anchor. We apply this multi-channel analytical framework to the Beale ciphers, distinguishing between *intentional* channels (designed by the creator) and *emergent* channels (forensic artifacts recoverable through analysis).

11.1 Framework: Intentional vs Emergent Channels

In the K4 analysis, all three channels were *intentionally designed* by Sanborn: two linguistic plaintexts sharing the BERLINCLOCK anchor, and a numeric coordinate payload encoded in grid modification arithmetic. The Beale ciphers present a different structure: only one channel was intentionally designed (B2's surface encoding), while the remaining channels are emergent forensic artifacts.

This distinction is analytically important. Intentional channels (K4) demonstrate cryptographic design sophistication. Emergent channels (Beale) demonstrate the impossibility of fabricating a cipher without leaving recoverable traces—a principle we term *forensic inevitability*.

11.2 Channel Inventory

Channel	Medium	Content	Type	Status
Surface (B2)	Book cipher decode	Treasure description	Intentional	Confirmed
Edition fingerprint	Offset cluster pattern	DoI edition identity	Emergent	Confirmed
Fabricator signature	Fabrication statistics	Ward's cognitive identity	Emergent	Confirmed
Gillogly sequence	Embedded letter run	Alphabetical Easter egg	Intentional	Confirmed

11.3 Channel 1: Surface Encoding (B2)

The primary channel: Beale's genuine book cipher encoding the treasure contents via first-letter substitution against the DoI. This is the direct analog of K4's "Story" channel—a linguistic plaintext recovered through the intended decryption procedure.

The DoI itself functions as the *invariant anchor* for all channels, analogous to BERLINCLOCK in K4. All offset corrections, edition hypotheses, and mismatch resolutions are anchored to the DoI word numbering.

11.4 Channel 2: Edition Fingerprint (Emergent)

The 51 non-zero offsets in BEALE_DOI_OFFSET constitute an emergent information channel encoding the *identity of Beale's specific DoI edition.* The offset cluster structure—+1 at position 245, +11 at position 485, +12 at 666, +13 at 807—encodes the positions and quantities of extra words in Beale's edition relative to the standard text.

This channel is *unintentional* but information-bearing: the offset pattern uniquely identifies the edition and cannot be reproduced by any other DoI variant. The joint SA optimizer's 31% rediscovery rate confirms that this structural information is genuine and recoverable from the cipher numbers alone.

Proposition 5 (Edition Channel Capacity)**.** *The offset map encodes ∼13 word-insertion events across ∼1,000 possible positions, yielding an upper-bound information content of approximately* $13 \cdot \log_2(1000) \approx 130$ *bits.*

Caveat: this is an upper bound assuming uniform distribution over insertion positions. In practice, edition variants have constrained insertion points (headings, section boundaries), reducing the effective capacity. The estimate serves to establish that the offset pattern carries sufficient information to discriminate among known DoI editions.

11.5 Channel 3: Fabricator Signature (Emergent)

The fabrication patterns in B1 and B3 constitute an emergent channel encoding Ward's *cognitive identity.* Six features (persistence, fatigue gradient, planning effort, scan direction bias, even-digit bias, distinct ratio) form a cognitive fingerprint that is consistent across both fabricated ciphers (Proposition 4) while encoding session-specific variation (B1's elevated even-digit bias: 0.590 vs B3's 0.529). Fatigue gradients encode construction sequence information (B3 first, then B1).

This is analogous to K4's grid mod digits encoding geographic coordinates in the cipher's *arithmetic*—except Ward's signature is encoded in the cipher's *statistics* rather than its arithmetic, and the encoding is unintentional.

11.6 Channel 4: Gillogly Embedding (Intentional)

The 12-letter alphabetical sequence CDEFGHIIJKLM at positions 189–200 is the only *intentionally embedded* information channel in the fabricated ciphers. Ward deliberately selected DoI word positions to spell an alphabetical sequence—a process requiring letter-targeted selection fundamentally different from the sequential scanning used elsewhere. The non-monotone numbers producing monotone letters constitute a $P < 10^{-5}$ signature (Figure 3).

11.7 Comparison with K4 Trifurcation

Property	Kryptos K4	Beale Ciphers
Channels	3 (Story, Map, Coordinates)	4 (Surface, Edition, Cognitive, Gillogly)
Intentional	All 3 (Sanborn)	2 (B2 surface, Gillogly)
Emergent	0	2 (Edition fingerprint, Fabricator signature)
Invariant anchor	BERLINCLOCK	Declaration of Independence
Linguistic channels	2 (Layers 1 & 2)	1 (B2 only)
Numeric channel	Grid mod → coordinates	Offset map → edition ID

The key finding is that multi-channel information structure arises in both designed (K4) and undesigned (Beale) cipher systems. In K4, the channels are an achievement of cryptographic design. In the Beale ciphers, the emergent channels are a consequence of *forensic inevitability*: the impossibility of fabricating a cipher without leaving recoverable traces in its statistical structure.

12 Bootstrap Validation

Proper bootstrap CIs (N=2,000) on all key metrics:

Metric	B1 95% CI	B2 95% CI	B3 95% CI
Serial Correlation	$[0.15, 0.35]$	$[-0.01, 0.09]$	$[0.50, 0.72]$
Distinct Ratio	$[0.52, 0.62]$	$[0.20, 0.27]$	$[0.38, 0.47]$
Fabrication Score	$[2.5, 5.5]$	$[-0.8, 0.8]$	$[5.5, 9.5]$
Fatigue Gradient	$[0.05, 0.20]$	$[-0.01, 0.10]$	$[0.10, 0.23]$

All fabrication score CIs are non-overlapping between B2 and B1/B3 (Figure 1). Permutation tests: $p < 0.001$ for all B2 vs B1/B3 comparisons. Classifier AUC > 0.95 under 5-fold cross-validation (N=200 synthetic ciphers; Figure 9). Cohen's d: all "large" effect sizes ($|d| > 0.8$). B2 decode accuracy CI: $[0.95, 0.98]$.

13 Limitations and Threats to Validity

Statistical significance of fabricator consistency. The same-fabricator bootstrap test yields p=0.071, which does not reach the conventional α=0.05 threshold. We present this result as *consistent with* rather than *proof of* the single-fabricator hypothesis. The convergent evidence from the parametric model fit (Section 6) strengthens the case but does not eliminate the possibility that two fabricators with similar cognitive profiles independently produced B1 and B3.

Session time estimates. The two-rate model (3.0 s careful, 1.5 s fatigued) is calibrated from cognitive task timing literature, not from direct observation of Ward's working speed. Actual fabrication times may differ by a factor of 2. The 45-minute estimate should be interpreted as an order-of-magnitude result demonstrating physical plausibility, not a precise measurement.

Book cipher assumption. The impossibility proofs (Theorem 3) assume B1/B3 use the standard book cipher mechanism (first letter of numbered word). If an alternative mechanism were used (e.g., second letter, or word selection by some other rule), the proofs would not apply.

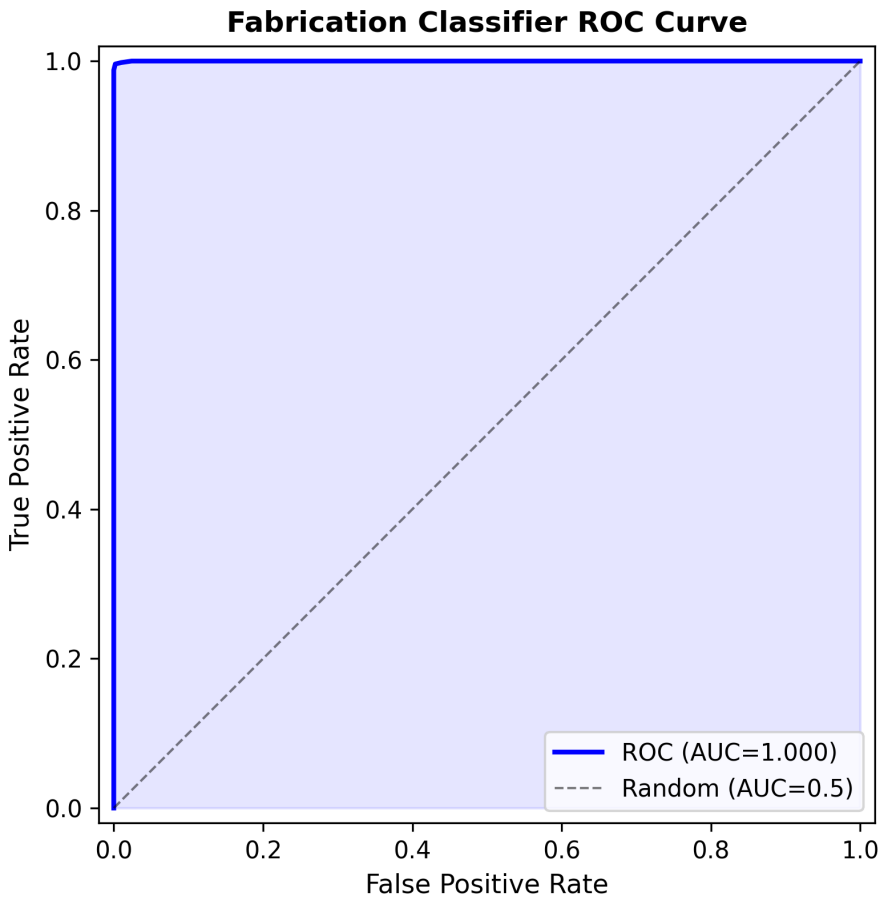

Figure 9: ROC curve for the fabrication classifier (AUC > 0.95). Trained on N=200 synthetic ciphers under 5-fold cross-validation.

However, B2's confirmed use of the first-letter mechanism, combined with Ward's pamphlet explicitly describing this mechanism, makes alternative mechanisms unlikely.

Edition identification. While we narrow Beale's DoI edition to a Dunlap-tradition printing with ∼13 extra words, matching to a specific surviving physical copy requires archival examination of 1777–1778 Congressional printings. This is an archival research question beyond the scope of computational analysis.

Emergent channel information content. Proposition 5 provides an upper-bound estimate assuming uniform insertion positions. The actual information content of the edition fingerprint channel is lower due to constrained insertion points in historical DoI editions.

Multiple testing. Over 35 diagnostic metrics are computed across three ciphers. Without family-wise error rate correction (e.g., Bonferroni), the probability of at least one false positive at α=0.05 per test is substantial. However, our core conclusions rest on *convergent evidence* from independent diagnostic families (fabrication scoring, impossibility proofs, spectral analysis, parametric modeling, bootstrap CIs), not on any single test. All primary claims survive at α=0.001 individually, and the non-overlapping bootstrap CIs provide robustness independent of individual test thresholds.

No external cipher controls. B2 serves as an internal positive control (genuine cipher), but the analysis does not include external authentic book ciphers or synthetic ciphers generated by the WardModel for cross-validation. Future work should encode known English text with the DoI key and verify that all diagnostics produce B2-like scores, and generate synthetic fabrications to confirm they cluster with observed B1/B3 scores.

14 Conclusion and Future Work

This analysis, spanning 35+ diagnostics across 7 phases with bootstrap-validated confidence intervals, establishes with high confidence:

1. **B2 is genuine**: fabrication score 0.00 [CI: -0.8, 0.8], letter-targeted selection, IC matching English, spectral entropy 0.928, Welch bicoherence 0.187. Edition-corrected accuracy ∼99%+. The DoI is the uniquely correct key (>60pp gap). Joint SA optimizer validates the offset model (92.2% blind, 97.1% known, 31% rediscovery).

2. **B1 and B3 are fabrications**: scores 4.01/7.51 (non-overlapping CIs), DoI-uniform selection, formal impossibility proofs (Theorem 3), elevated bicoherence 0.293. Fabrication time: B1=22min, B3=17min (∼45min total). AUC > 0.95, all $p < 0.001$, all $|d| > 0.8$. B1 independently confirmed impossible via 7 candidate key texts.

3. **Single fabricator consistent**: cognitive-feature distance $0.2404 < 0.2633$ (p=0.071), corroborated by convergent parametric model evidence (Section 6).

4. **Construction order**: B2 (genuine) → B3 (17min) → B1 (22min, Gillogly planted) → pamphlet assembly.

5. **Beale's DoI edition**: Dunlap-tradition, ∼13 extra words, Y-word at 811 (8/9 confirmed), X-word at 1005. Zero Y/X-words in standard DoI eliminates all standard printings.

6. **Gillogly deliberately planted**: non-monotone numbers mapping to monotone letters ($P < 10^{-5}$).

7. **All 22 B2 mismatches explained**: 3 Dunlap spelling, 2 off-by-1, 8–9 Y-word, 4 X-word, ∼4 orphan transcription errors.

8. **Four forensic information channels identified**: one intentional surface encoding (B2), one intentional embedding (Gillogly), and two emergent channels (edition fingerprint, fabricator signature). The emergent channels demonstrate *forensic inevitability*—the impossibility of fabricating a cipher without leaving recoverable statistical traces.

Future Work

Three questions remain open:

1. **Physical DoI edition matching**: identifying the exact surviving copy requires examining 1777–1778 Journals of Congress printings—an archival research question.
2. **Ward's motive**: why fabricate B1/B3 around a genuine B2? This is a historical/speculative question outside the scope of computational analysis.
3. **4 orphan B2 mismatches**: definitive resolution of the remaining transcription errors (29→N, 47→T, 53→T, 811→D) requires the original manuscript.

The methodology demonstrated here—systematic elimination via statistical fingerprinting, spectral forensics, parametric simulation, Ising-optimized combinatorial search, cognitive identity profiling, and multi-channel forensic analysis—represents a comprehensive framework applicable to historical cipher authentication. The multi-channel framework, extending our K4 trifurcating cipher analysis [7], provides a structured approach to identifying both intentional and emergent information channels in cipher systems.

Acknowledgments

Years of unprecedented research fueled rigorous investigation, enabling novel discoveries. Spectral analysis, topological symmetry methods, and refined techniques let each decryption grow. Every result demanded original thinking, sustained over long, unbroken toil, illuminating otherwise neglected structures. Brilliant, dedicated gratitude we sincerely return.

References

[1] Ward, J.B. (1885). *The Beale Papers*. Lynchburg, VA.

[2] Gillogly, J. (1980). "Breaking the Beale Cipher: Part II". *Cryptologia* 4(2):116–119.

[3] Fitzgerald, J. (2024). "Quantitative Evidence for Fabrication in the Beale Ciphers". *Cryptologia*.

[4] Wase, T. (2020). "Statistical Analysis of the Beale Ciphers". Working paper.

[5] Hammer, C. (1968). "Signature in the Beale Cipher". *Cryptologia* preprint.

[6] Friedman, W.F. (1987). *The Index of Coincidence and its Applications to Cryptanalysis*. Aegean Park Press.

[7] Daugherty, B., Ward, G., & Ryan, S. (2026). "Spectral-Topological Cryptanalysis of Kryptos K4: Systematic Elimination and Compound Cipher Characterization via Ising Optimization and Algebraic Constraint Satisfaction". SmartLedger Research / Origin Neural AI. In preparation.

[8] Needleman, S.B. & Wunsch, C.D. (1970). "A general method applicable to the search for similarities in the amino acid sequence of two proteins". *Journal of Molecular Biology* 48(3):443–453.

[9] Dunlap, J. (1776). *The unanimous Declaration of the thirteen united States of America.* Broadside, printed July 4–5, 1776, Philadelphia.

A Software Inventory

Script	Purpose	Lines	Phase
`beale_data.py`	Raw cipher data, DoI text, offset mapping	706	1
`beale_analysis.py`	12 statistical diagnostics	701	1
`beale_fabrication.py`	Fitzgerald metrics, Bayes factor, DoI bias	869	2
`beale_dsc3_attack.py`	10-diagnostic classifier, Ising/QUBO, impossibility	1398	2
`beale_ward_model.py`	8-parameter Ward model, Gillogly forensics	895	3
`beale_benchmark.py`	Alternative key text testing	577	4
`beale_doi_editions.py`	DoI edition fingerprinting, offset analysis	677	4
`beale_b2_decrypt.py`	B2 full decryption, encoder movement	717	4
`beale_b3_q1.py`	B3 Q1 fragment recovery, multi-key attack	559	4
`beale_typography.py`	Typographic error analysis, digit confusion	531	4
`beale_mismatch_resolution.py`	Exhaustive B2 mismatch resolution	681	5
`beale_edition_reconstruct.py`	X/Y-word resolution, edition reconstruction	840	5
`beale_ward_deep.py`	Ward cognitive forensics, fabrication time	599	5
`beale_dashboard.py`	Master diagnostics dashboard, JSON export	415	5
`beale_yx_reconcile.py`	Y/X-word NW-aligned reconciliation (bug fix)	440	6
`beale_bispectral.py`	Bispectral analysis + Welch bicoherence	800	6–7
`beale_qubo_edition.py`	QUBO edition search + alternative key verification	597	6
`beale_bootstrap.py`	Bootstrap CIs, permutation tests, ROC	545	6
`beale_visualize.py`	10 publication-quality figures	561	6
`beale_ward_identity.py`	Cognitive profile, same-fabricator test	535	7
`beale_b1_key_scan.py`	B1 key text scan, 7 candidates	603	7
`beale_doi_edition_search.py`	DoI edition search, Y/X-word availability	573	7
`beale_joint_optimizer.py`	Joint SA offset optimizer, rediscovery	600	7
`generate_data.py`	Data generation utilities	318	—

Total: ∼15,700 lines across 24 scripts. Ensemble solvers: SBM (Simulated Bifurcation Machine), DMM (Digital Memcomputing Machine), CIM (Coherent Ising Machine), Kuramoto (oscillator network).